

# Tools and Strategies for Tracing Memory Usage and Leaks

Justin Farnsworth, Matt Mozdzen, Matthew Van Soelen, Brian Worts

To Contact for further information on this topic: [farnswj1@tcnj.edu](mailto:farnswj1@tcnj.edu),  
[mozdzen1@tcnj.edu](mailto:mozdzen1@tcnj.edu), [vansoem1@tcnj.edu](mailto:vansoem1@tcnj.edu), and [wortsb1@tcnj.edu](mailto:wortsb1@tcnj.edu)

Canvas Page: <https://tcnj.instructure.com/groups/364604>

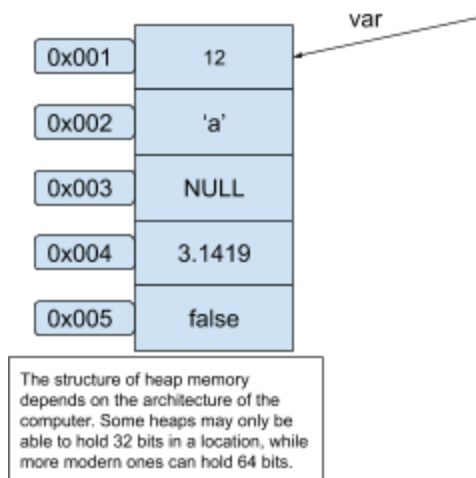
## What is a Pointer?

A *pointer* is a reference to a memory location of a variable. In C, you can declare a pointer variable like so:

```
int x = 12;
```

```
int *var = &x;
```

//var would have the value "0x001"



## Pointers in Ruby

In Ruby, variables are pointers to objects. However, Ruby can support the usage of C-type pointers with the class `Fiddle::Pointer`.

You can declare a pointer as such:

```
Fiddle::Pointer.new(address) → fiddle_cptr
```

```
#declare a c-type pointer
```

## What is Memory Tracing?


Memory tracing is the process of monitoring and recording any changes in a system's memory. This makes it a very effective tool for debugging. Most approaches involve executing additional code every time a memory address is accessed. As a result, memory tracing can increase performance overhead by up to 10 times in some cases, which is why it is usually not enabled by default.

## What is a Memory Leak?



A **memory leak** is an event that occurs when memory is allocated for an object/variable but continues to take up that memory after it is no longer needed by a running program, thus resulting in decreased available heap memory for other processes. A very common sign of a memory leak is that the longer an application is running, the more memory gets consumed. If a memory leak occurs, they can negatively impact performance and even cause failure/crashing of the application.

## Causes



Memory leaks are caused by poor programming and fall into many categories. In general, they result from a program not properly deallocating blocks of memory while running or even after being closed.

One such reason for a memory leak is a "*strong reference cycle*." This means that objects contain references to each other, thus preventing deallocation of the memory.

Another such reason is that, as a program runs a block of code repeatedly, instead of allocating that space and then reallocating that space for the next iteration, it allocates a new chunk of memory each time.

## **Solutions for Memory Leaks**

On the user end, one of the simplest ways to resolve problems caused by a memory leak is to close the application causing the leak or, if that does not work, then restarting their computer will. Restarting the computer “clears” the memory of the computer by deallocating used-up space. However, running the program that caused the memory leak again will cause another buildup in the memory and the computer will need to be restarted again. Thus, the only way to stop a memory leak in a program is for the developer to remove it from the code, which is why memory tracing is such an important skill for developers to know.

## **Commonly Used Tools for Memory Tracing**

### **GDB (GNU Debugger)**

Useful for debugging C and C++ programs, mainly for identifying problems involving memory leaks. To debug memory leaks with GDB, you will need to use the built-in tool Valgrind.

Step 1: Launch Valgrind

```
1 valgrind --vgdb=yes --vgdb-error=0 <program> <arguments>
```

STEP 2: Follow the printed instructions to run the gdb

```
1 ==21399== TO DEBUG THIS PROCESS USING GDB: start GDB like this
2 ==21399== /path/to/gdb &lt;path to program>;
3 ==21399== and then give GDB the following command
4 ==21399== target remote | /usr/lib/valgrind/../../bin/vgdb --pid=21399
```

STEP 3: Launch the gdb for your program

```
1 gdb <program>
```

STEP 4: Disable Non-Stop Mode

```
1 gdb> set non-stop off
```

STEP 5: Attach Valgrind to your gdb

```
1 gdb> target remote | vgdb
```

STEP 6: Check for memory leaks

```
1 gdb> monitor leak_check
```

Now, run the program and leak check at an occasional interval. Stop it when a leak is found. Restart the program and do this, making the interval smaller each time, until you can find where the leak is coming from. If you have an idea of where the leak is, put a break statement after it in your code.

GDB Manual: <https://sourceware.org/gdb/current/onlinedocs/gdb/>

Source:

<http://www.responsive.se/thomas/2013/09/20/debugging-memory-leaks-with-valgrind-and-gdb/>

## Garbage Collectors

Some languages, such as JAVA, have a built-in program that automatically manages memory. This makes cleaning up memory more convenient the programmer doesn't have to manually clean it up.

However, it doesn't necessarily mean that garbage collectors are faster than manual memory management. If the programmer knows what needs to be cleaned up, then it may be efficient to manually clean it up, as the garbage collector will allocate possibly more time and resources to do so. Some garbage collectors will even pause the program while it cleans up memory.

## **Task Manager**

Although minimal detail is given, the "Task Manager" (on Windows) can be used to monitor RAM usage (along with some other hardware component usage). This is useful when a user is running many applications but does not know which one is slowing down their computer's performance.

## **Memory Tracing with XCode**


### **Memory Graph Debugger**

This is a tool provided by XCode to help developers find memory leaks for XCode versions 8.0 and up. It visualizes all the objects in memory while the program is running, and then guides the user to where that memory leak is occurring.

In order to visualize the memory use, this tool provides a helpful graph in real time of the memory being used while the program is running. A memory leak can be clearly seen when, for example, the user hits a button over and over again and each time, the memory use of the program increases.

Then, to help find the leak, objects stored in the memory are shown. The user can scan this list and should they find multiple instances of the same object in the heap, it becomes obvious what is causing the leak.

Memory  
Duration: 34 sec  
High: 92,7 MB  
Low: Zero KB

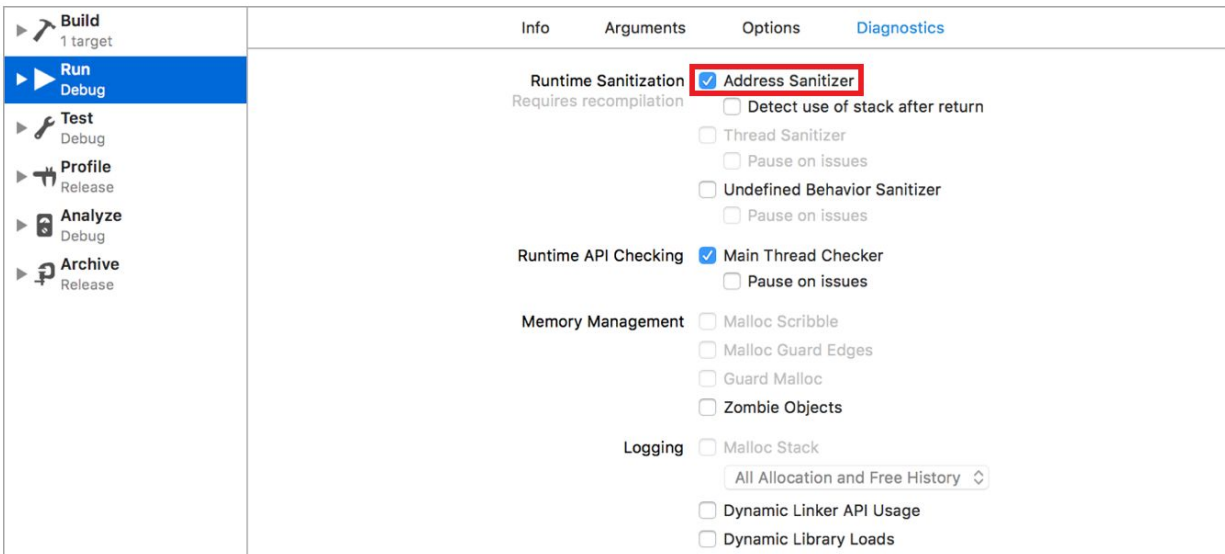


The Memory Graph Debugger tool showing memory usage over time of an application repeatedly running the same function with a memory leak in it.

Source:  
<https://medium.com/@imas145/using-memory-graph-debugger-in-xcode-9-a50aa9d0243f>

Examples of how to use this tool are given in Hands On Examples 1 and 2.

Another useful tool in XCode is the **Address Sanitizer**, which adds instrumentation to your app that enables Xcode to stop your app where memory corruption happens. Address sanitizer finds problems such as accessing deallocated pointers, buffer overflow and underflow of the heap and stack, and other memory issues. To use address sanitizer, enable it in the debug scheme for your target, then run and use the app. Xcode monitors memory use and stops your app on the line of code causing the problem and opens the debugger. Use the debugger to isolate the cause.



Running your code with **Address Sanitizer** checks enabled typically results in CPU slowdown of

2× to 5×, and an increase in memory usage by 2× to 3×. You can improve memory utilization by compiling at the -O(1) optimization level.

For most use cases, the overhead introduced by the Address Sanitizer should be acceptable for daily use during development—in fact, you may not even notice any slowdown.

## **Further Reading**

<https://useyourloaf.com/blog/xcode-visual-memory-debugger/>

<https://medium.com/@imas145/using-memory-graph-debugger-in-xcode-9-a50aa9d0243f>

<https://agostini.tech/2018/12/09/memory-graph-debugging-in-xcode/>

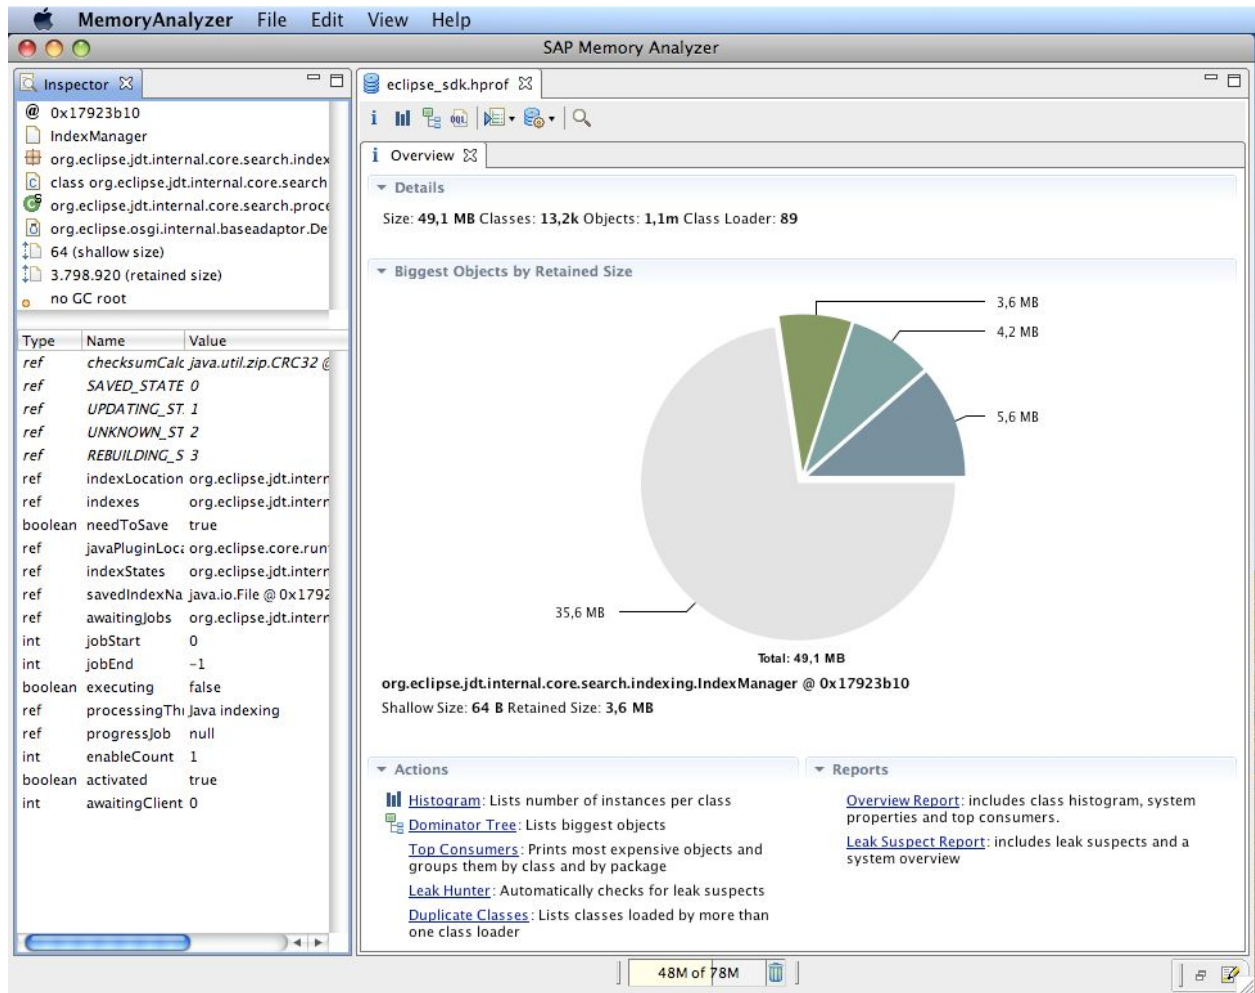
<https://www.surekhatech.com/blog/find-memory-leaks-in-ios-apps-with-xcode-instruments>

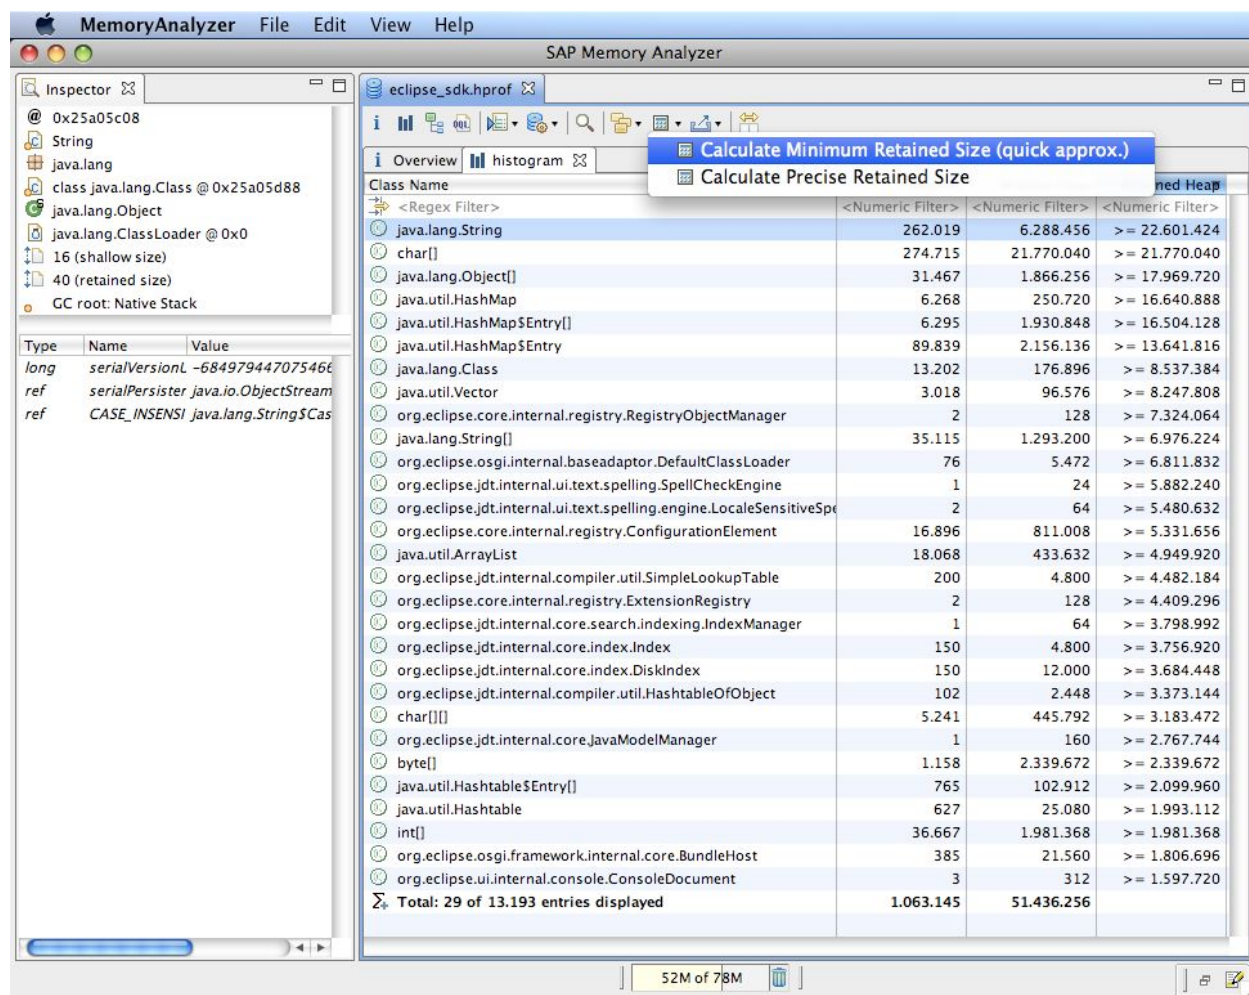
[https://developer.apple.com/documentation/code\\_diagnostics/address\\_sanitizer](https://developer.apple.com/documentation/code_diagnostics/address_sanitizer)

## **Eclipse Memory Analyzer (MAT)**

Eclipse allows you to monitor memory usage and find memory leaks. It is particularly useful for Java heap-dumps. Conveniently, MAT visualizes memory usage across programs through various diagrams, such as pie charts and histograms.







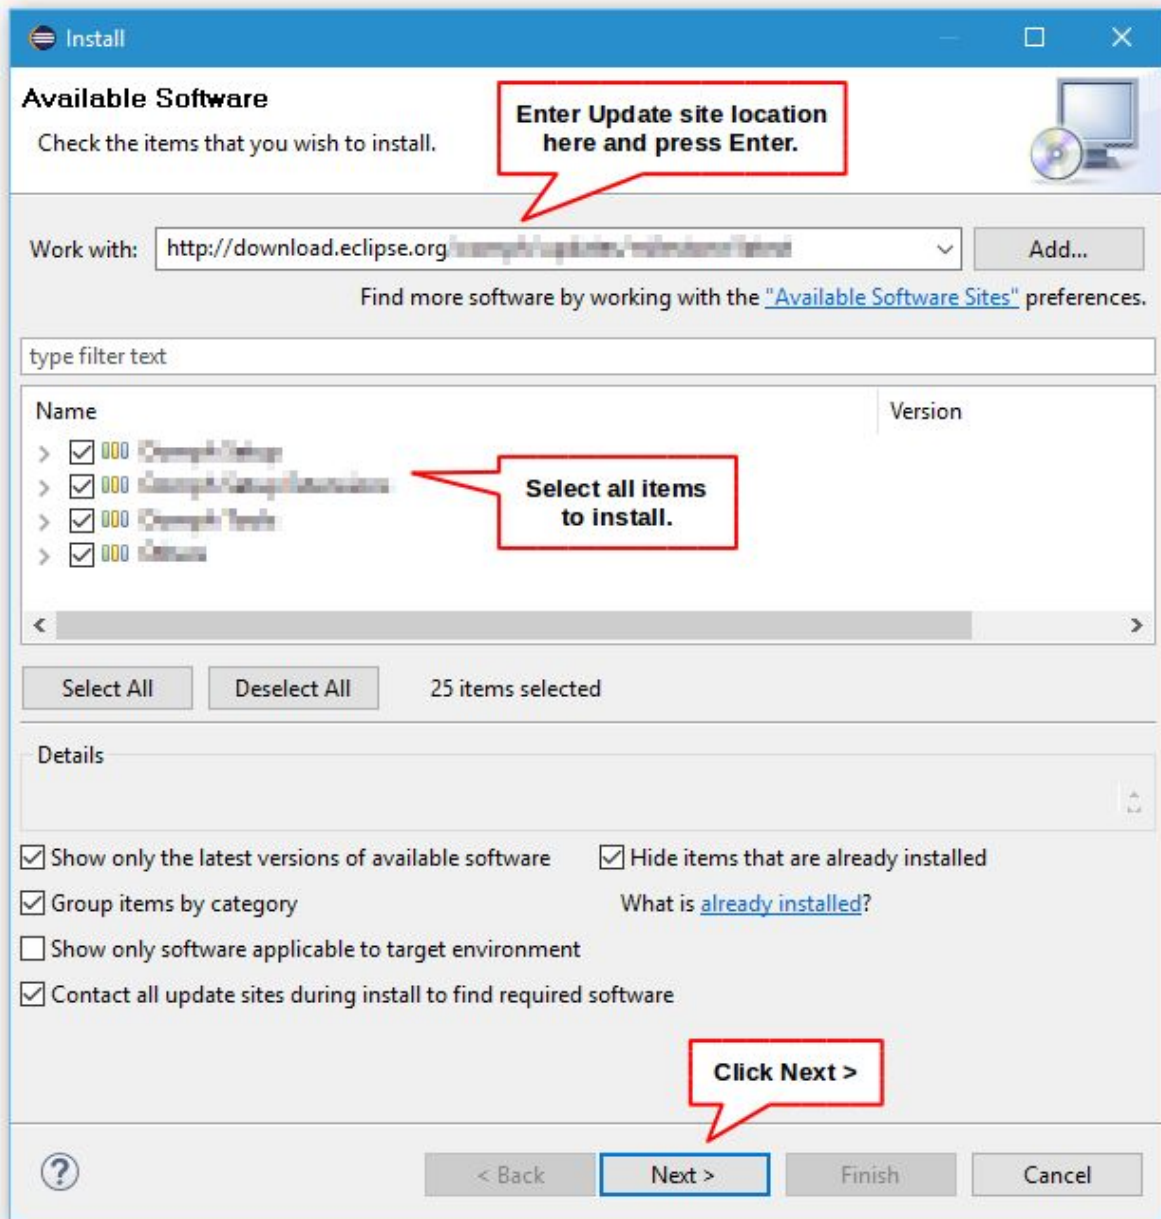
## Installation

To install MAT, first ensure that the Eclipse IDE is installed.

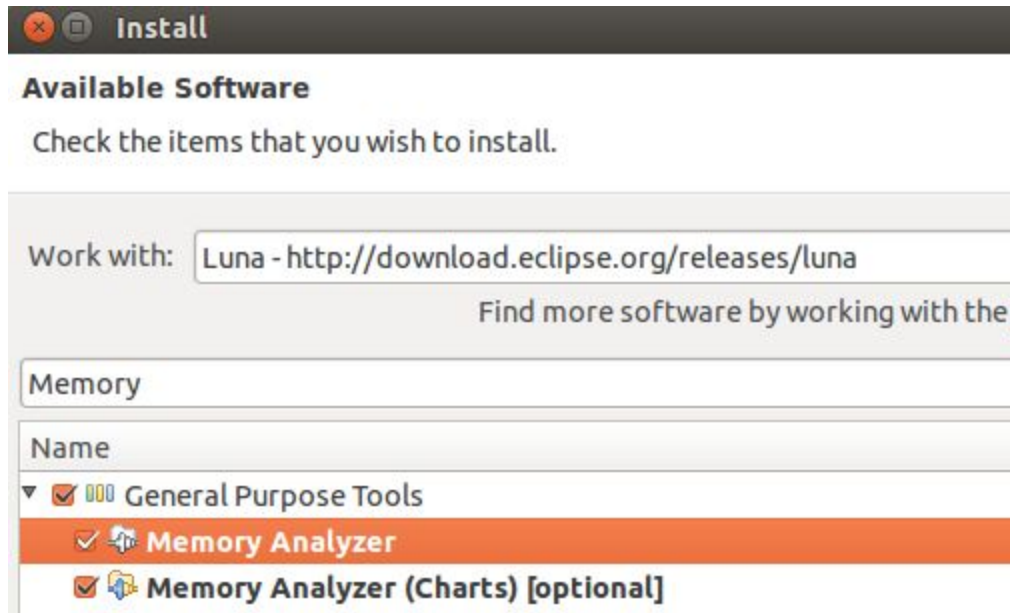
Then go to Help > Install New Software...

In "Work with", insert the update link and download the contents.

E.g.: <https://download.eclipse.org/mat/1.4/update-site/>



Select *General Purpose Tools*, *Memory Analyzer* and *Memory Analyzer (Charts)*.



## Utilizing MAT

First, create a heap dump by running the application in Eclipse with the argument:

`-XX:+HeapDumpOnOutOfMemoryError`

Go to File > New > Other > Heap Dump to set up the heap dump.

Once the heap dump has been created (look for the .hprof ending), open it via double-clicking it.

Open your project and click Leak Suspects Report. Then click Finish.

From there, you can analyze the heap dump and the objects that are used in the program. The dominator tree provides information on the used objects, making it useful to identify memory leaks.

## Further Reading

<https://www.eclipse.org/mat/>

<https://www.vogella.com/tutorials/EclipseMemoryAnalyzer/article.html#installation>

# Memory Analysis Over Time in Ruby

## REQUIRED RUBY FILES TO INSTALL:

createdb.rb

db.rb

gemfile

gencsv.rb

genimport.sh

Graph.rb

All of these files can be found at the Github link for this method.

You can use the Ruby command *ObjectSpace* to obtain information regarding objects that may be sticking around. This can be used to dump all objects in memory to a file at spaced time intervals (in the source's example, an interval of 5 minutes is used). The objects that are consistently being dumped to the file are the ones that are causing memory leakage.



To implement the object dump in your program, first you need to include the gem *objspace*. Based upon whether or not you are developing a Rails application, there are two ways to create a dump function. If you are not creating a Rails application, follow Example 1. If you are, follow Example 2.

## Example 1 - (Regular Ruby Application File):

NOTE: This method assumes that you have a main event loop that you can call every X minutes you want to dump the data.

Create a function called `heap_dump` that mimics the following code closely:

```
def heap_dump
  GC.start

  l = Time.now.strftime('%s')

  open("/tmp/ruby-heap-#{l}.dump", "w") do |io|
    objectSpace.dump_all(output: io)
  end
end
```

Before you main function code, insert the following:

```
ObjectSpace.trace_object_allocations_start
```

At the end of your main loop, call the `heap_dump` function.

For Example:

```
mainloop do
  # put your code here

  heap_dump
end
```

## Example 2 - (For Rails Applications):

Add the following code to a controller in your Rails application that you visit every X minutes (the amount of time you want to track between object dumps):

```
class HeapDumpsController < ActionController::Metal

  def heap_dump

    if ENV['HEAP_DUMP'] == '1' && params[:token].to_s == ENV['HEAP_DUMP_TOKEN']

      heap_dump

      self.response_body = 'Dumped heap'

    else

      self.status = 401

      self.response_body = 'Invalid token'

    end

  end

end
```

Next, add the following line to your configuration file:

```
get "/heap_dump", to: HeapDumpsController.action(:heap_dump)
```

Finally, add the next code snippet to your heap dump tracing file:

```
if ENV['HEAP_DUMP'] == 1

  require 'objspace'

  ObjectSpace.trace_object_allocations_start

end
```

Now, to analyze the memory leak, take the heap\_dump text file you generated and load it into a database. First, initialize the database using the following command in the terminal:

**bundle exec ruby createdb.rb**

Next, import your data. Importing the data will take some time, so the importing process is split into two steps; conversion to SQL files, and then loading the converted data into the database. Enter the following on the command line to do so:


```
bundle exec ruby gencsv.rb
```

```
sh genimport.sh | psql mem_analysis
```

Once all of the data is imported, then we need to create the graph using the following terminal command:

```
bundle exec ruby graph.rb type-mem
```

A graph will be generated with all object types as a line graph each, plotted as Time vs. # of Object Type Count. The types that have a flat slope are not the troublemakers. The object type that has an increasing slope is the one that is causing the memory leak.



If you want to look further into which variable exactly is causing the issue, create another graph of just that object type using:



**bundle exec ruby graph.rb [object type]-mem**

If you still cannot find the issue, you may need to look into modifying the graphs.rb file to include more cases of how the memory leak could occur.

## **Further Reading**

<https://gist.github.com/wvengen/f1097651c238b2f7f11d>

## **Hands-On Exercises**

Link to first exercise: <https://tcnj.instructure.com/courses/1630707/pages/group-6-activity-1>

Link to second exercise: <https://tcnj.instructure.com/courses/1630707/pages/group-6-activity-2>